

AKSHAT GUPTA

CS PhD, UC Berkeley

 [Personal Website](#)  akshat.gupta@berkeley.edu
 412 892 0560  [Linkedin](#)
 Berkeley, California  [Google Scholar](#)

SUMMARY

I'm a second year PhD Student at UC Berkeley affiliated with BAIR. I'm currently working on interpretability with a focus on creating more efficient and interpretable AI.

SKILLS

Languages: Python, Go, MATLAB, C++, Javascript, C

Tools: Pytorch, Huggingface, NLTK, Spacy

Technologies: AWS, GCP, Azure, Dialogflow, Docker

EDUCATION

8/2023 - Present	PhD, Computer Science Advisor : Gopala Anumanchipalli	University of California, Berkeley
1/2020 - 5/2021	MS, Electrical and Computer Engineering Advisor : Alan W Black	Carnegie Mellon University
10/2017 - 3/2020	MS, Applied and Engineering Physics Thesis Advisor : Gregory Eyink (Johns Hopkins University)	Technical University of Munich, Germany
8/2013 - 6/2017	B.Tech, Electrical Engineering Thesis Advisor : Renu Rameshan	Indian Institute of Technology Mandi

EXPERIENCE

6/2021 - 8/2023	Senior Associate, AI Research <ul style="list-style-type: none">Developed relation extraction models for financial knowledge graphs, resulting in a SIGIR 2023 paper [9] and a patent [3]Led the design and deployment of solutions for public information discovery, enabling insights into companies and financial institutions, resulting in a patent [1]Supervised a summer intern and mentored student teams from Columbia University and Carnegie Mellon University on sponsored capstone projects, fostering collaborative research and development	JPMorgan Chase
9/2018 - 6/2019	Visiting Research Scholar <ul style="list-style-type: none">Implemented stochastic differential equations (SDE) to study to run large-scale turbulence simulation on Hopkins Cluster (MARCC)Investigated resultant vorticity dynamics in turbulent channel flows under Lagrangian DynmaicsResulted in 2 publications [13, 14] in Journal of Fluid Mechanics, the #1 journal in the field	Johns Hopkins

PUBLICATIONS

- How Linearly Associative are Memories in Large Language Models?**
Akshat Gupta, Nehal Sindhu, Gopala Anumanchipalli
Oral Presentation, ICLR 2025 Workshop (New Frontiers in Associative Memories)
- Norm Growth and Stability Challenges in Localized Sequential Knowledge Editing**
Akshat Gupta, Christine Fang, Atahan Ozdemir, Maochuan Lu, Ahmed Alaa, Thomas Hartvigsen, Gopala Anumanchipalli
KnowFM @ AAAI 2025 (Outstanding Paper Award)
- Sylber: Syllabic Embedding Representation of Speech from Raw Audio**
Cheol Jun Cho, Nicholas Lee, **Akshat Gupta**, Dhruv Agarwal, Ethan Chen, Alan Black, Gopala Anumanchipalli
ICLR 2025
- PokerBench : Training Large Language Models to become Professional Poker Players**
Richard Zhuang, **Akshat Gupta**, Richard Yang, Aniket Rahane, Zhengyu Li, Gopala Anumanchipalli
AAAI 2025
- Rebuilding ROME : Resolving Model Collapse during Sequential Model Editing**
Akshat Gupta, Sidharth Baskaran, Gopala Anumanchipalli
EMNLP 2024
- A Unified Framework for Model Editing**
Akshat Gupta, Dev Sajnani, Gopala Anumanchipalli
EMNLP 2024 Findings
- Model Editing at Scale leads to Gradual and Catastrophic Forgetting**
Akshat Gupta, Anurag Rao, Gopala Anumanchipalli
ACL 2024 Findings

8. **Self-Assessment Tests are Unreliable Measures of LLM Personality**
Akshat Gupta, Xiaoyang Song, Gopala Anumanchipalli
BlackboxNLP 2024, co-located with EMNLP 2024
9. **REFinD: Relation Extraction Financial Dataset**
Simerjeet Kaur, Charese Smiley, Akshat Gupta, Joy Sain, Dongsheng Wang, Toyin Aguda, Sameena Shah
SIGIR 2023
10. **Probing Quantifier Comprehension in Large Language Models: Another Example of Inverse Scaling**
Akshat Gupta
BlackboxNLP 2023, co-located with EMNLP 2023
11. **Intent classification using pre-trained language agnostic embeddings for low resource languages**
Hemant Yadav, Akshat Gupta, Sai Krishna Rallabandi, Alan W Black, Rajiv Ratn Shah
Interspeech 2022
12. **Acoustics Based Intent Recognition Using Discovered Phonetic Units for Low Resource Languages**
Akshat Gupta, Sai Krishna Rallabandi, Alan W Black
ICASSP 2021
13. **Stochastic Lagrangian dynamics of vorticity. Part 1. General theory for viscous, incompressible fluids**
Gregory Eyink, Akshat Gupta, Tamer Zaki
Journal of Fluid Mechanics, 2020 (Volume 901, Page A2)
14. **Stochastic Lagrangian dynamics of vorticity. Part 2. Application to near-wall channel-flow turbulence**
Gregory Eyink, Akshat Gupta, Tamer Zaki
Journal of Fluid Mechanics, 2020 (Volume 901, Page A3)

PATENTS

1. **Method and system for automated public information discovery**
Akshat Gupta, Simerjot Kaur, Xiaomo Liu, Armineh Nourbakhsh, Andrea Stefanucci, Alex Woodgate, Sameena Shah
US Patent 12,001,491
2. **Method and system for detection of anomalous rejections of foreign exchange requests**
Nacho Navarro, Xiaomo Liu, Simran Lamba, Akshat Gupta, Sameena Shah
US Patent App. 18/085,104
3. **Method and system for understanding financial documents**
Simerjot Kaur, Charese Smiley, Joy Sain, Suchetha Siddagangappa, Akshat Gupta, Sameena Shah
US Patent App. 17/647,356

AWARDS

- Outstanding Paper Award, KnowFM Workshop @ AAAI 2025 [2]
- NVIDIA Compute Grant of 16,000 GPU hours for A100 GPUs (80GB) on the topic - "Enhancing Post-training Editability via Knowledge Modularization"
- Student Travel Grant, BlackboxNLP 2024, co-located with EMNLP 2024
- UC Berkeley EECS Fellowship, 2023
- Best Undergraduate Thesis, IIT Mandi 2017 [Topic - Blind Image Deconvolution]

ACADEMIC SERVICE

Reviewing

- ICLR 2025, ICML 2025, COLM 2025
- ACL Rolling Review : June 2024, August 2024, October 2024, December 2024
- ACL 2023
- EMNLP 2022

Workshop Organization

- 4th Workshop on Knowledge Discovery from Unstructured Data in Financial Services, SIGIR 2023

ADVISING

- | | |
|---|----------------|
| • Maochuan Lu, Undergraduate UC Berkeley, Published [2] | 2024 - Present |
| • Richard Zhuang, Undergraduate UC Berkeley, Published [4] | 2024 - Present |
| • Dev Sajjani, Undergraduate UC Berkeley, Published [6] | 2024 |
| • Anurag Rao, Undergraduate UC Berkeley (→ MS, University of Oxford), Published [7] | 2023 - 2024 |
| • Xiaoyang Song, MS Columbia University (→ PhD, University of Michigan), Published [8] | 2023 |
| • Anant Singh, MS NYU (→ Machine Learning Engineer, Apple) | 2022 - 2023 |